

Clase 7

Diccionario de datos:

https://www.nyc.gov/assets/tlc/downloads/pdf/data_dictionary_trip_records_yellow.pdf

1. En Hive, crear la siguiente tabla (externa) en la base de datos tripdata:

airport_trips(tpep_pickup_datetime, airport_fee, payment_type, tolls_amount, total_amount)

```
CREATE EXTERNAL TABLE airport_trips (tpep_pickup_datetime date, airport_fee float,
payment_type int, tolls_amount float, total_amount float)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/tables/external/airport_trips';
```

2. En Hive, mostrar el esquema de airport_trips

```
hive> describe formatted airport_trips;
OK
# col_name          data_type          comment
tpep_pickup_datetime  date
airport_fee           float
payment_type          int
tolls_amount          float
total_amount          float

# Detailed Table Information
Database:              tripdb
Owner:                 hadoop
CreateTime:            Tue May 14 10:39:25 ART 2024
LastAccessTime:        UNKNOWN
Retention:             0
Location:              hdfs://172.17.0.2:9000/tables/external/airport_trips
Table Type:            EXTERNAL_TABLE
Table Parameters:
    EXTERNAL           TRUE
    numFiles            18
    totalSize           626116905
    transient_lastDdlTime 1716224111

# Storage Information
SerDe Library:         org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:           org.apache.hadoop.mapred.TextInputFormat
OutputFormat:          org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:            No
Num Buckets:           -1
Bucket Columns:        []
Sort Columns:          []
Storage Desc Params:
    field.delim         ,
    serialization.format ,
Time taken: 0.292 seconds, Fetched: 33 row(s)
```

3. Crear un archivo .bash que permita descargar los archivos mencionados abajo e ingestarlos en HDFS:

```
wget -O /home/hadoop/landing/Yellow_tripdata_2021-01.parquet
https://edvaibucket.blob.core.windows.net/data-engineer-edvai/yellow_tripdata_2021-01.parquet?
sp\=r\&st\=2023-11-06T12:52:39Z\&se\=2025-11-06T20:52:39Z\&sv\=2022-11-02\&sr\=c\&sig\
=J4Ddi2c7Ep23OhQLPisbYaerlH472iigPwc1%2FkG80EM%3D
```

```
/home/hadoop/hadoop/bin/hdfs dfs -put -f /home/hadoop/landing/Yellow_tripdata_2021-01.parquet
/ingest/airflow
```

```
hadoop@401bec58e4c6:~/airflow/dags$ hdfs dfs -ls /ingest/airflow
Found 2 items
-rw-r--r--  1 hadoop supergroup    21686067 2024-05-20 13:53 /ingest/airflow/Yellow_tripdata_2021-01.
parquet
-rw-r--r--  1 hadoop supergroup    21777258 2024-05-20 13:54 /ingest/airflow/Yellow_tripdata_2021-02.
parquet
```

4. Crear un archivo .py que permita, mediante Spark, crear un data frame uniendo los viajes del mes 01 y mes 02 del año 2021 y luego Insertar en la tabla airport_trips los viajes que tuvieron como inicio o destino aeropuertos, que hayan pagado con dinero.

```
from pyspark.context import SparkContext
from pyspark.sql.session import SparkSession
from pyspark.sql import HiveContext

sc = SparkContext('local')
spark = SparkSession(sc)
hc = HiveContext(sc)

df01 = spark.read.option("header", "true").parquet("hdfs://172.17.0.2:9000/ingest/airflow/
Yellow_tripdata_2021-01.parquet")

df02 = spark.read.option("header", "true").parquet("hdfs://172.17.0.2:9000/ingest/airflow/
Yellow_tripdata_2021-02.parquet")

total_trips_df = df01.unionAll(df02)

total_trips_filtered_df = total_trips_df.filter("RatecodeID = 2 AND payment_type = 2")

final_total_trips_df =
total_trips_filtered_df.select(total_trips_filtered_df.tpep_pickup_datetime.cast("date"),
total_trips_filtered_df.airport_fee.cast("float"), total_trips_filtered_df.payment_type.cast("int"),
total_trips_filtered_df.tolls_amount.cast("float"), total_trips_filtered_df.total_amount.cast("float"))

#final_total_trips_df.show(10)

final_total_trips_df.createOrReplaceTempView("airport_trips_view")

hc.sql("insert into tripdb.airport_trips select * from airport_trips_view;")
```

5. Realizar un proceso automático en Airflow que orqueste los archivos creados en los puntos 3 y 4. Correrlo y mostrar una captura de pantalla (del DAG y del resultado en la base de datos)

```
from datetime import timedelta
from airflow import DAG
from airflow.operators.bash import BashOperator
```

```

from airflow.operators.dummy import DummyOperator
from airflow.utils.dates import days_ago

args = {
    'owner': 'airflow',
}

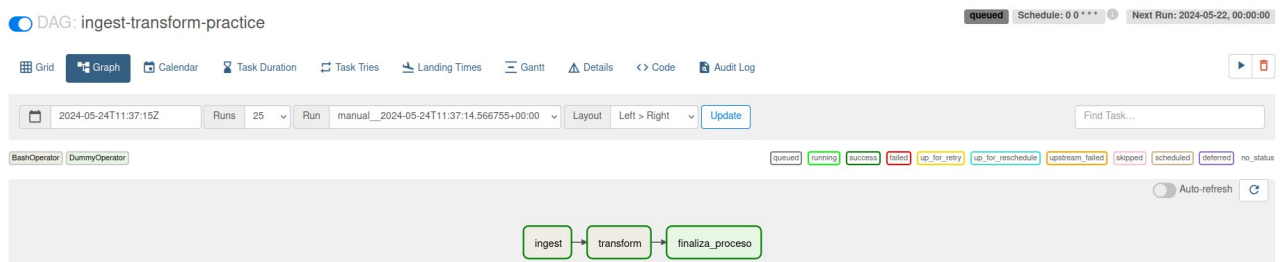
with DAG(
    dag_id='ingest-transform-practice',
    default_args=args,
    schedule_interval='0 0 * * *',
    start_date=days_ago(2),
    dagrun_timeout=timedelta(minutes=60),
    tags=['ingest', 'transform'],
    params={"example_key": "example_value"},
) as dag:

    finaliza_proceso = DummyOperator(
        task_id='finaliza_proceso',
    )
    ingest = BashOperator(
        task_id='ingest',
        bash_command='/usr/bin/sh /home/hadoop/scripts/practica_airflow/ingest.sh ',
    )
    transform = BashOperator(
        task_id='transform',
        bash_command='ssh hadoop@172.17.0.2 /home/hadoop/spark/bin/spark-submit --files
/home/hadoop/hive/conf/hive-site.xml /home/hadoop/scripts/practica_airflow/transformation.py ',
    )

    ingest >> transform >> finaliza_proceso

if __name__ == "__main__":
    dag.cli()

```



Luego de ejecutarse el dag, en hive ya se han cargado los datos a la tabla airport_trips

```

hive> show databases;
OK
default
f1
tripdb
Time taken: 0.891 seconds, Fetched: 3 row(s)

```

```
hive> use tripdb;
OK
Time taken: 0.041 seconds
```

```
hive> show tables;
OK
airport_trips
congestion
distance
passenger
payments
tolls
Time taken: 0.088 seconds, Fetched: 6 row(s)
```

```
hive> select * from airport_trips limit 10;
```

```
hive> select * from airport_trips limit 10;
OK
2020-12-31      NULL      2          0.0        11.8
2020-12-31      NULL      2          0.0         4.3
2020-12-31      NULL      1          0.0        51.95
2020-12-31      NULL      1          0.0        36.35
2020-12-31      NULL      1          0.0        24.36
2020-12-31      NULL      1          0.0        14.15
2020-12-31      NULL      2          0.0        17.3
2020-12-31      NULL      2          0.0        21.8
2020-12-31      NULL      4          0.0        28.8
2020-12-31      NULL      1          0.0        18.95
Time taken: 2.012 seconds, Fetched: 10 row(s)
```

Clase 8

Diccionario de datos:

<https://www.kaggle.com/datasets/rohanrao/formula-1-world-championship-1950-2020?select=results.csv>

1. Crear la siguientes tablas externas en la base de datos f1 en hive:

- driver_results (driver_forename, driver_surname, driver_nationality, points)
- constructor_results (constructorRef, cons_name, cons_nationality, url, points)

```
create external table driver_results(  
  driver_forename string,  
  driver_surname string,  
  driver_nationality string,  
  points int)  
COMMENT 'Driver Results'  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
LOCATION '/tables/external/f1/';
```

```
create external table constructor_results(  
  constructor_ref string,  
  cons_name string,  
  cons_nationality string,  
  url string,  
  points int)  
comment 'Constructor Results'  
row format delimited  
fields terminated by ','  
location '/tables/external/f1/';
```

2. En Hive, mostrar el esquema de driver_results y constructor_results

```
hive> describe formatted driver_results;  
OK  
# col_name      data_type      comment  
forename        string  
surname         string  
nationality     string  
points          bigint  
  
# Detailed Table Information  
Database:      f1  
Owner:         hadoop  
CreateTime:    Fri May 24 08:19:47 ART 2024  
LastAccessTime: UNKNOWN  
Retention:     0  
Location:      hdfs://172.17.0.2:9000/user/hive/warehouse/f1.db/driver_results  
Table Type:    MANAGED_TABLE  
Table Parameters:  
  numFiles              1  
  spark.sql.create.version  3.2.0  
  spark.sql.sources.provider  parquet  
  spark.sql.sources.schema  {\"type\":\"struct\", \"fields\": [{\"name\":\"forename\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"surname\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"nationality\", \"type\":\"string\", \"nullable\":true, \"metadata\":{}}, {\"name\":\"points\", \"type\":\"long\", \"nullable\":true, \"metadata\":{}}]}  
  totalSize             14565  
  transient_lastDdlTime  1716549587  
  
# Storage Information  
SerDe Library:  org.apache.hadoop.hive ql.io.parquet.serde.ParquetHiveSerDe  
InputFormat:    org.apache.hadoop.hive ql.io.parquet.MapredParquetInputFormat  
OutputFormat:   org.apache.hadoop.hive ql.io.parquet.MapredParquetOutputFormat  
Compressed:     No  
Num Buckets:    -1  
Bucket Columns: []  
Sort Columns:   []  
Storage Desc Params:  
  path              hdfs://172.17.0.2:9000/user/hive/warehouse/f1.db/driver_results  
  serialization.format  1  
Time taken: 0.144 seconds, Fetched: 34 row(s)
```

```

hive> describe formatted constructor_results;
OK
# col_name          data_type          comment
constructor_ref      string
cons_name            string
cons_nationality     string
url                  string
points               bigint

# Detailed Table Information
Database:            f1
Owner:               hadoop
CreateTime:          Fri May 24 08:19:51 ART 2024
LastAccessTime:      UNKNOWN
Retention:           0
Location:            hdfs://172.17.0.2:9000/user/hive/warehouse/f1.db/constructor_results
Table Type:          MANAGED_TABLE
Table Parameters:
    numFiles          1
    spark.sql.create.version 3.2.0
    spark.sql.sources.provider  parquet
    spark.sql.sources.schema [{"type": "struct", "fields": [{"name": "constructor_ref", "type": "string", "nullable": true, "metadata": {}}, {"name": "cons_name", "type": "string", "nullable": true, "metadata": {}}, {"name": "cons_nationality", "type": "string", "nullable": true, "metadata": {}}, {"name": "url", "type": "string", "nullable": true, "metadata": {}}, {"name": "points", "type": "long", "nullable": true, "metadata": {}}]}]
    nullability        true
    totalSize          2601
    transient_lastDdlTime 1716549591

# Storage Information
Serde Library:       org.apache.hadoop.hive ql.parquet.serde.ParquetHiveSerde
InputFormat:         org.apache.hadoop.hive ql.parquet.MapredParquetInputFormat
OutputFormat:        org.apache.hadoop.hive ql.parquet.MapredParquetOutputFormat
Compressed:          No
Num Buckets:         -1
Bucket Columns:      []
Sort Columns:        []
Storage Desc Params:
    path              hdfs://172.17.0.2:9000/user/hive/warehouse/f1.db/constructor_results
    serialization.format 1
Time taken: 0.163 seconds, Fetched: 35 row(s)

```

3. Crear un archivo .bash que permita descargar los archivos mencionados

```
wget -O /home/hadoop/landing/results.csv https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/results.csv
```

```
hdfs dfs -put /home/hadoop/landing/results.csv /ingest/f1
```

```
wget -O /home/hadoop/landing/drivers.csv https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/drivers.csv
```

```
hdfs dfs -put /home/hadoop/landing/drivers.csv /ingest/f1
```

```
wget -O /home/hadoop/landing/constructors.csv
```

```
https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/constructors.csv
```

```
hdfs dfs -put /home/hadoop/landing/constructors.csv /ingest/f1
```

```
wget -O /home/hadoop/landing/races.csv https://dataengineerpublic.blob.core.windows.net/data-engineer/f1/races.csv
```

```
hdfs dfs -put /home/hadoop/landing/races.csv /ingest/f1
```

4. Generar un archivo .py que permita, mediante Spark:

a. insertar en la tabla driver_results los corredores con mayor cantidad de puntos en la historia.

b. insertar en la tabla constructor_result quienes obtuvieron más puntos en el Spanish Grand Prix en el año 1991

```
from pyspark.sql.session import SparkSession
```

```

spark = SparkSession.builder \
    .appName("CSV to Hive") \
    .config("spark.sql.catalog.Implementation", "hive") \
    .enableHiveSupport() \
    .getOrCreate()

```

```
df_results = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/f1/results.csv")
```

```
df_drivers = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/f1/drivers.csv")
```

```
df_constructors = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/f1/constructors.csv")
```

```
df_races = spark.read.option("header", "true").csv("hdfs://172.17.0.2:9000/ingest/f1/races.csv")
```

```

df_results.createOrReplaceTempView("vw_results")
df_drivers.createOrReplaceTempView("vw_drivers")
df_constructors.createOrReplaceTempView("vw_constructors")
df_races.createOrReplaceTempView("vw_races")

df_driver_results = spark.sql( \
    "SELECT \
        cast(d.forename as string), \
        cast(d.surname as string), \
        cast(d.nationality as string), \
        sum(cast(r.points as int)) points \
    FROM \
        vw_results r, \
        vw_drivers d \
    WHERE \
        r.driverId = d.driverId \
    GROUP BY \
        d.driverId, d.forename, d.surname, d.nationality \
    ORDER BY \
        points DESC" \
    )

df_constructor_results = spark.sql( \
    "SELECT \
        cast(c.constructorRef as string) constructor_ref, \
        cast(c.name as string) cons_name, \
        cast(c.nationality as string) cons_nationality, \
        cast(c.url as string) url, \
        sum(cast(r.points as int)) points \
    FROM \
        vw_results r, \
        vw_constructors c, \
        vw_races ra \
    WHERE \
        r.constructorId = c.constructorId \
        AND r.raceId = ra.raceId \
        AND ra.year = '1991' \
    GROUP BY \
        c.constructorId, constructor_ref, cons_name, cons_nationality, c.url \
    ORDER BY \
        points DESC" \
    )

df_driver_results.write.mode("overwrite").saveAsTable("f1.driver_results")
df_constructor_results.write.mode("overwrite").saveAsTable("f1.constructor_results")

```

5. Realizar un proceso automático en Airflow que orqueste los archivos creados en los puntos 3 y 4. Correrlo y mostrar una captura de pantalla (del DAG y del resultado en la base de datos)

```

from datetime import timedelta
from airflow import DAG
from airflow.operators.bash import BashOperator
from airflow.operators.dummy import DummyOperator
from airflow.utils.dates import days_ago

args = {
    'owner': 'airflow',
}

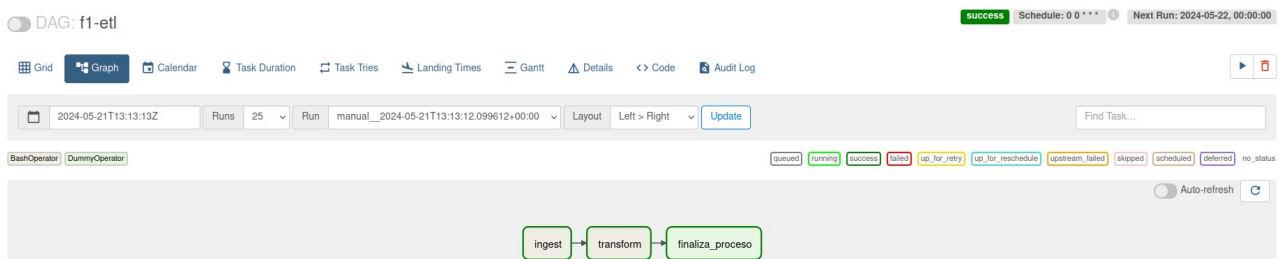
with DAG(
    dag_id='f1-etl',
    default_args=args,
    schedule_interval='0 0 * * *',
    start_date=days_ago(2),
    dagrun_timeout=timedelta(minutes=60),
    tags=['ingest', 'transform'],
    params={"example_key": "example_value"},
) as dag:

    finaliza_proceso = DummyOperator(
        task_id='finaliza_proceso',
    )
    ingest = BashOperator(
        task_id='ingest',
        bash_command='/usr/bin/sh
/home/hadoop/scripts/practica_airflow/clase8/ingest.sh ',
    )
    transform = BashOperator(
        task_id='transform',
        bash_command='ssh hadoop@172.17.0.2 /home/hadoop/spark/bin/spark-submit
--files /home/hadoop/hive/conf/hive-site.xml
/home/hadoop/scripts/practica_airflow/f1-transformation2.py ',
    )

    ingest >> transform >> finaliza_proceso

if __name__ == "__main__":
    dag.cli()

```



```

hive> select * from constructor_results limit 10;
OK
mclaren McLaren British http://en.wikipedia.org/wiki/McLaren 139
williams Williams British http://en.wikipedia.org/wiki/Williams_Grand_Prix_Engineering
ferrari Ferrari Italian http://en.wikipedia.org/wiki/Scuderia_Ferrari 55
benetton Benetton Italian http://en.wikipedia.org/wiki/Benetton_Formula 38
jordan Jordan Irish http://en.wikipedia.org/wiki/Jordan_Grand_Prix 13
tyrrell Tyrrell British http://en.wikipedia.org/wiki/Tyrrell_Racing 12
minardi Minardi Italian http://en.wikipedia.org/wiki/Minardi 6
dallara Dallara Italian http://en.wikipedia.org/wiki/Dallara 5
brabham Brabham British http://en.wikipedia.org/wiki/Brabham 3
team_lotus Team Lotus British http://en.wikipedia.org/wiki/Team_Lotus 3
Time taken: 0.245 seconds, Fetched: 10 row(s)

```



```
hive> select * from driver_results limit 10;
OK
Lewis      Hamilton      British 4308
Sebastian   Vettel        German  3077
Fernando    Alonso        Spanish 2021
Kimi        Räikkönen     Finnish 1873
Max         Verstappen    Dutch   1792
Valtteri    Bottas        Finnish 1775
Nico        Rosberg       German  1594
Michael     Schumacher    German  1566
Daniel      Ricciardo     Australian 1289
Jenson      Button        British 1235
Time taken: 0.223 seconds, Fetched: 10 row(s)
```

BONUS: GroupTask

En lugar de la tarea de ingest anterior que contenía todas las sentencias para la descarga y puesta en hdfs de los archivos, se separó por dos task group donde en una se ejecutan los comandos para los archivos de conductores y constructores, mientras que en el otro grupo se ejecutan las tareas para los archivos de carreras y resultados.

```
with TaskGroup("ingest") as ingest:

    inicializa_ingest = EmptyOperator(task_id='inicializa_ingest',)

    finaliza_ingest = EmptyOperator(task_id='finaliza_ingest',)

    with TaskGroup("ingest_actors") as ingest_actors:

        ingestDrivers = BashOperator(
            task_id='ingest-drivers',
            bash_command='/usr/bin/sh/home/hadoop/scripts/practica_airflow/clase8/ingest-drivers.sh ',
        )

        ingestConstructors = BashOperator(
            task_id='ingest-constructors',
            bash_command='/usr/bin/sh/home/hadoop/scripts/practica_airflow/clase8/ingest-constructors.sh ',
        )

    with TaskGroup("ingest_metrics") as ingest_metrics:

        ingestResults = BashOperator(
            task_id='ingest-results',
            bash_command='/usr/bin/sh/home/hadoop/scripts/practica_airflow/clase8/ingest-results.sh ',
        )

        ingestRaces = BashOperator(
            task_id='ingest-races',
            bash_command='/usr/bin/sh/home/hadoop/scripts/practica_airflow/clase8/ingest-races.sh ',
        )

    inicializa_ingest >> [ingest_actors, ingest_metrics] >> finaliza_ingest
inicializa_proceso >> ingest >> transform >> finaliza_proceso
```

